

The landscape of Parallel Computing Research: A view from Berkeley

by Yunqi Zhang

04/05/2013

1. Applications

Trying to better express a parallel computation with 13 dwarfs.

2. Hardware

- a. Processors: building smaller processor building blocks
 - i. energy-efficient
 - ii. dynamic voltage scaling and power down
 - iii. heterogeneity
- b. Memory: much larger fraction of total system silicon for memory
- c. Interconnection networks: linearly scale
 - i. collective - small message size
 - ii. point-to-point - large message size
 - iii. exception: 3-D FFT
- d. Communication Primitives: Synchronization
 - i. transactional memory
 - ii. full-empty bits in memory
 - iii. message passing
- e. Performance and energy counters:
 - i. programmer and compiler are explicitly responsible for performance

3. Programming models

- a. psychological research (really???)
- b. independent of the number of processors -> scalability
- c. rich set of data sizes and types
- d. support of proven styles of parallelism

4. System software

- a. Autotuners
 - i. larger search space by parallelism
- b. Deconstructing the operating system support
 - i. embedded systems increase in functionality
 - ii. virtual machine: protection against attacks, failures don't interfere other programs

5. Metrics of success

- a. Maximizing programmer productivity
- b. maximizing application performance