

A Primer on Memory Consistency and Cache Coherence (Chapter 3-4)

by Yunqi Zhang

04/15/2013

1. Memory Consistency Model (MC Model)
 - a. give rules that partition executions into those obeying MC and those disobeying MC
 - b. compare to coherence:
 - i. coherence: per cache block basis
 - ii. consistency: across all blocks
2. Sequential Consistency (SC)
 - a. memory order respects each core's program order, where memory order is the order of operations
 - b. memory order respects each core's program order
 - c. naive implementations
 - i. uniprocessor: executing all threads on a single sequential core
 - ii. switch: switch selects to execute one operation from one of the multi-cores in its original program order
 - iii. (+) allow SC executions
 - iv. (-) does not scale up with increasing core count, sequential performance bottleneck
 - d. basic implementation with cache coherence
 - i. Single-Writer-Multiple-Reader invariant cache coherence
 - ii. each core has its own cache
 - iii. (+) operations can be done concurrently in parallel if they do not conflict
 - e. impacts from more features
 - i. non-binding prefetching: As the loads and stores are performed in program order, it doesn't matter in what order coherence permissions are obtained
 - ii. speculative cores: similar to non-binding prefetching that no effect on SC
 - iii. **A core can issue coherence requests in any order (don't have to be program order)**
 - iv. dynamically scheduling: check required
 1. speculatively accessed block has not left the cache
 2. replay each speculative load check if the value loaded is equal to previous one
 - v. non-binding prefetching in dynamically scheduling:
 1. **SC only requires a core's loads and stores access its level-one cache in program order**
 2. **SC dictates the order in which loads and stores get applied to coherent memory but does NOT dictate the order of coherence activity**

- vi. multithreading: need to ensure a thread cannot read a value written by another thread on the same core before the store has been made **VISIBLE**
 - f. atomic operations
 - i. atomically perform “read-modify-write”: “test-and-set, fetch-and-increment”
 - ii. spin-lock: use RMW to atomically read whether lock’s value is unlocked
 - g. MIPS R10000
 - i. issue loads and stores in program order into address queue
3. Total Store Order
- a. motivation: hide the latency of frequently happened stores
 - b. single-core: bypassing the value of the most recent store to A to succeeding loads
 - c. TSO
 - i. preserve the program order for: 1) load->load, 2) load->store, 3) store->load
 - ii. omit store->load (does not matter for most programs)
 - iii. programmers can prevent non-SC execution by adding FENCE for above
 - d. loads and stores leave each core in that core’s program order <p
 - e. a load either bypasses a value from the write buffer or awaits the switch as before
 - f. a store enters the tail of the FIFO write buffer or stalls the core if the buffer is full
 - g. when the switch selects core Ci, it performs either the next load or the store at the head of the write buffer
 - h. **A thread may bypass values that it has written, but other threads may not see the value until the store is inserted into the memory order**
 - i. FENCE: specify that all instructions before the FENCE in program order must be ordered before any instructions after the FENCE in program order