

A Primer on Memory Consistency and Cache Coherence (Chapter 6-7)

by Yunqi Zhang

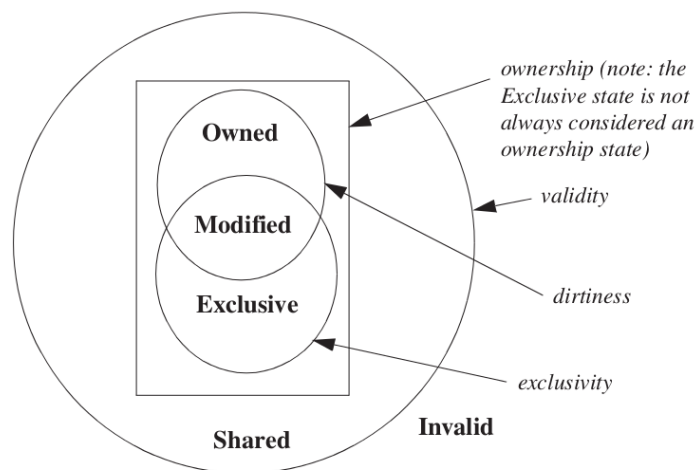
04/19/2013

1. Coherence Protocol

- a. each storage structure has a finite state machine called a **coherence controller**
- b. one independent, identical finite state machine per block
- c. cache controller
 - i. request from core side: loads and stores from the core, returns values back
 - 1. cache miss: coherence transaction start with a coherence request through the interconnection network
 - ii. interface to network side: connects to the rest of the system
- d. memory controller (LLC, Memory)
 - i. only network side interface: doesn't issue or receive coherence requests
- e. can be clearly described by tabular

2. Design Space

- a. state - things we want to encode
 - i. validity: most up-to-date value for the block
 - ii. dirtiness: most up-to-date value but differs from LLC/MEM
 - iii. exclusivity: only privately cached copy in the system, can be cached in LLC
 - iv. ownership: responsible for responding to coherence requests
- b. state - one commonly used **stable** states
 - i. M(odified): valid, exclusive, owned, and potentially dirty - RW
 - ii. S(hared): valid, ! exclusive, ! owned, ! dirty - R
 - iii. I(nvalid): ! valid - NULL
 - iv. O(wned): valid, ! exclusive, owned, and potentially dirty - R
 - v. E(xclusive): valid, exclusive, ! dirty - R



- c. state - **transient** states between the transition of two stable states
 - d. state - blocks in LLC/MEM
 - i. cache-centric (common): aggregation of the states of that block in the caches
 - ii. memory-centric: corresponds to the memory controller's permissions to that block
 - e. state - maintain the block state
 - i. cache: few bits for each block
 - ii. memory:
 - 1. inclusive LLC maintains a copy of every block that is cached somewhere
- its state in memory is the same as in LLC
 - f. Transactions
 - i. GetShared: obtain block in Shared (read-only) state
 - ii. GetModified: obtain block in Modified (read-only) state
 - iii. Upgrade: upgrade block state from read-only to read-write
 - iv. PutShared: evict block in Shared state
 - v. PutExclusive: evict block in Exclusive state
 - vi. PutOwned: evict block in Owned state
 - vii. PutModified: evict block in Modified state
3. Major Protocol Design Options
- a. Snooping VS Directory
 - i. snooping: broadcast a request message to all other coherence controllers
 - 1. (+) simple
 - 2. (-) do not scale to large numbers of cores
 - ii. directory: unicast a request message to the controller that is the home of that block
 - 1. (+) scalability
 - 2. (-) more time or message for many transactions
 - b. Invalidate VS Update
 - i. invalidate: invalidate the copies in other caches when a core wish to write
 - ii. update: update the copies in other caches when a core wish to write
 - 1. (+) reduce the latency for a core to read a newly written block
 - 2. (-) consume more bandwidth
 - 3. (-) complicate the implementation of memory consistency model